

Below screenshot shows content of table emp-

```
hive (employees)> select * from emp;
OK
emp.emp_id      emp.emp_name    emp.sal emp.dept
101      Amitabh 20000      1
102      Shahrukh 10000      2
103      Akshay 11000      3
104      Anubhav 5000       4
105      Pawan 2500       5
106      Aamir 25000      1
107      Salman 17500      2
108      Ranbir 14000      3
109      Katrina 1000       4
110      Priyanka 2000       5
111      Tushar 500        1
112      Ajay 5000       2
113      Jubeen 1000       1
114      Madhuri 2000       2
Time taken: 4.364 seconds, Fetched: 14 row(s)
hive (employees)>
```

Below query was used to find list of all employees who draw higher salary than the average salary of that department-

```
SELECT temp.emp_name, temp.sal, temp.dept, temp.avg_salary
FROM
( SELECT avg(sal) OVER (PARTITION BY dept) AS avg_salary, emp_id, emp_name, sal, dept FROM emp ) temp
WHERE temp.sal > temp.avg_salary;
```

Below screenshot shows the query and results for it-

```
hive (employees)> SELECT temp.emp_name, temp.sal, temp.dept, temp.avg_salary FROM
> ( SELECT avg(sal) OVER (PARTITION BY dept) AS avg_salary, emp_id, emp_name, sal, dept FROM emp ) temp
> WHERE temp.sal > temp.avg_salary;
Query ID = acadgild_20171114165959_6c7ef08e-5700-4541-bad2-a6c6c0a8b971
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1510658819198_0001, Tracking URL = http://localhost:8088/proxy/application_1510658819198_0001/
Kill Command = /home/acadgild/hadoop-2.6.0/bin/hadoop job -kill job_1510658819198_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2017-11-14 17:01:13,484 Stage-1 map = 0%, reduce = 0%
2017-11-14 17:01:57,125 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 13.58 sec
2017-11-14 17:02:27,227 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 21.44 sec
2017-11-14 17:02:34,012 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 32.56 sec
MapReduce Total cumulative CPU time: 32 seconds 560 msec
Ended Job = job_1510658819198_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 32.56 sec HDFS Read: 506 HDFS Write: 157 SUCCESS
Total MapReduce CPU Time Spent: 32 seconds 560 msec
OK
temp.emp_name    temp.sal      temp.dept      temp.avg_salary
Aamir 25000      1      11625.0
Amitabh 20000      1      11625.0
Shahrukh 10000      2      8625.0
Salman 17500      2      8625.0
Ranbir 14000      3      12500.0
Anubhav 5000       4      3000.0
Pawan 2500       5      2250.0
Time taken: 167.756 seconds, Fetched: 7 row(s)
hive (employees)>
```